

# Double ML & Heterogeneous Treatment Effects

Assignment 4

Zaheer Abbas

2025-07-02

## Table of contents

1	Loading Dataset	1
2	ATE of Notification Sent via Double ML (Partial-Linear Model)	2
3	ATE of Notification Sent via Double ML (AIPW Model)	3
4	Verification with DoubleML package	4
5	Heterogeneous Treatment Effects of Notification Sent	6
6	HTE Evaluation	8
7	Optimal Policy Learning	10

## 1 Loading Dataset

```
print(dataset.head(6))
```

	age	gender	region	sessions_per_day	past_engagement	device_type	\
0	54	Male	Asia	3	34.831284	iOS	
1	18	Male	US	5	77.758422	iOS	
2	42	Male	US	1	53.025246	iOS	
3	27	Female	US	2	19.752821	Android	
4	53	Female	US	1	79.309742	iOS	
5	35	Female	EU	7	71.962521	Android	

	app_version	notification_sent	pop_up_rem	notification_opened	engagement
0	1	1	0	0	14.924706
1	2	0	0	0	5.229459
2	1	1	1	1	27.129607
3	1	0	1	1	17.782530
4	2	0	0	1	25.818239
5	1	1	1	0	23.795170

## 2 ATE of Notification Sent via Double ML (Partial-Linear Model)

```
# Final OLS regression: y_resid ~ t_resid
model = sm.OLS(y_resid_all, t_resid_all_with_const)
results = model.fit()

# Print the ATE estimate
print("Estimated ATE (DML):", results.params[1])
print("Standard Error:", results.bse[1])
print(results.summary())
```

Estimated ATE (DML): 8.739531269869078

Standard Error: 0.16458490838508583

### OLS Regression Results

```
=====
Dep. Variable:          y      R-squared:          0.220
Model:                  OLS    Adj. R-squared:      0.220
Method:                 Least Squares    F-statistic:      2820.
Date:                   Wed, 02 Jul 2025    Prob (F-statistic): 0.00
Time:                   14:46:24    Log-Likelihood:    -34722.
```

No. Observations:	10000	AIC:	6.945e+04
Df Residuals:	9998	BIC:	6.946e+04
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	0.0192	0.078	0.247	0.805	-0.134	0.172
x1	8.7395	0.165	53.100	0.000	8.417	9.062
Omnibus:	7.281	Durbin-Watson:	2.010			
Prob(Omnibus):	0.026	Jarque-Bera (JB):	6.782			
Skew:	0.030	Prob(JB):	0.0337			
Kurtosis:	2.887	Cond. No.	2.11			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly speci

### 3 ATE of Notification Sent via Double ML (AIPW Model)

```
# Compute AIPW scores
aipw_scores = (
    T * (Y - mu1_preds) / e_preds + mu1_preds
    - (1 - T) * (Y - mu0_preds) / (1 - e_preds) - mu0_preds
)

# Estimate ATE and standard error
ate = np.mean(aipw_scores)
se = np.std(aipw_scores) / np.sqrt(len(aipw_scores))

print(f"Estimated ATE (AIPW): {ate:.4f}")
print(f"Standard Error: {se:.4f}")
```

Estimated ATE (AIPW): 8.6743  
Standard Error: 0.2725

## 4 Verification with DoubleML package

```
# Create and fit DoubleMLPLR model
dml_plr_obj = dml.DoubleMLPLR(obj_dml_data, ml_l, ml_m)
dml_plr_obj.fit()

# Print results
print(dml_plr_obj.summary)
```

	coef	std err	t	P> t	2.5 %	97.5 %
notification_sent	8.764015	0.166246	52.717287	0.0	8.43818	9.08985


```
# DoubleMLIRM model
dml_irm_obj = dml.DoubleMLIRM(obj_dml_data, ml_g=ml_g_cloned, ml_m=ml_m_cloned)
dml_irm_obj.fit()

# Print results
print(dml_irm_obj.summary)
```

	coef	std err	t	P> t	2.5 %	97.5 %
notification_sent	8.722585	0.164558	53.006115	0.0	8.400057	9.045113

```
methods= ["Manual PLM", "Manual AIPW", "DoubleML PLR", "DoubleML IRW"]
estimates = [manual_plm_ate, manual_aipw_ate, dml_plr_ate, dml_irm_ate]

#plot
plt.figure(figsize=(8,5))
plt.bar(methods, estimates, color=['orange', 'green', 'blue', 'purple'])
plt.ylabel('ATE Estimate')
plt.title('Comparison of ATE Estimates from Different Methods')
plt.show()
```



#### 4.0.1 Interpretation,

All four methods yield very similar ATE estimates, around 8.7 to 8.9, indicating consistency across both manual and automated (DoubleML-based) implementations. This suggests that the estimated treatment effect is robust to the choice of method, providing confidence in the reliability of the causal inference results. # ATE of Notification Opened Via DoubleML

```
dml_iivm_obj = dml.DoubleMLIIVM(obj_dml_data, ml_g, ml_m, ml_r)
dml_iivm_obj.fit()

ate_estimate = dml_iivm_obj.coef[0]
# Display
print("\nATE Estimate:", ate_estimate)

print(dml_iivm_obj.summary)
```

ATE Estimate: 8.39439199546623

	coef	std err	t	P> t	2.5 %	\
notification_opened	8.394392	0.847088	9.909707	3.777532e-23	6.73413	

	97.5 %
notification_opened	10.054654

## 5 Heterogeneous Treatment Effects of Notification Sent

- Estimate group ATEs by gender (M/F) and by tertiles of past\_engagement.

```
print("\nGATE by Gender:")
print(gate_gender)
print("\nGATE by Past Engagement:")
print(gate_engagement)
```

GATE by Gender:

===== DoubleMLBLP Object =====

----- Fit summary -----

	coef	std err	t	P> t	[0.025	\
Group_Female	10.053458	0.246148	40.843166	0.000000e+00	9.571017	
Group_Male	7.358292	0.215963	34.072013	1.916407e-254	6.935012	

	0.975]
Group_Female	10.535899
Group_Male	7.781571

GATE by Past Engagement:

===== DoubleMLBLP Object =====

----- Fit summary -----

	coef	std err	t	P> t	[0.025	\
Group_Low	7.324093	0.301686	24.277244	3.410174e-130	6.732800	
Group_Medium	8.448609	0.271020	31.173416	2.443172e-213	7.917420	
Group_High	10.394889	0.278911	37.269595	5.102941e-304	9.848234	

	0.975]
Group_Low	7.915385
Group_Medium	8.979797
Group_High	10.941543

- R-Learner

RandomForestRegressor()

```
ite_rl = tau_model.predict(X_eval)
eval_data['ite_rl'] = ite_rl
print(eval_data[['ite_rl']].head())
```

	ite_rl
6252	5.848759
4684	6.275323
1731	0.974028
4742	4.754528
4521	2.290479

- DR-Learner model

```
X_eval = pd.get_dummies(eval_data[X_cols], drop_first=True)
X_eval = X_eval.reindex(columns=X_train.columns, fill_value=0)

ite_dr_eval = tau_model.predict(X_eval) # tau_model is DR-Learner model
eval_data['ite_dr'] = ite_dr_eval

print(eval_data[['ite_dr']].head())
```

	ite_dr
6252	5.848759
4684	6.275323
1731	0.974028
4742	4.754528
4521	2.290479

## 6 HTE Evaluation


LinearRegression()

```
print("BLP coefficient:", blp_model.coef_[0])
print("BLP intercept:", blp_model.intercept_)

plt.scatter(ite_dr_eval, aipw_scores, alpha=0.3)
plt.plot(ite_dr_eval, blp_model.predict(ite_dr_eval.reshape(-1,1)), color='red')
plt.xlabel("Predicted ITE (DR-Learner)")
plt.ylabel("AIPW Score")
plt.title("Best Linear Predictor (BLP) Plot")
plt.show()
```

BLP coefficient: 0.9843571937512781

BLP intercept: 7.944396402807552



```
print("Samples per group:")
print(counts)
```



```

print("\nSorted Group Average Treatment Effects (GATES):")
print(gates)

# Plotting the GATES
plt.figure(figsize=(8, 5))
plt.bar(gates.index.astype(str), gates.values, color='skyblue')
plt.xlabel('Group (Sorted by Predicted ITE)')
plt.ylabel('Average Treatment Effect (ITE - Doubly Robust)')
plt.title('Sorted Group Average Treatment Effect (GATES)')
plt.show()

```

Samples per group:

```

group
1      400
2      400
3      400
4      400
5      400
Name: count, dtype: int64

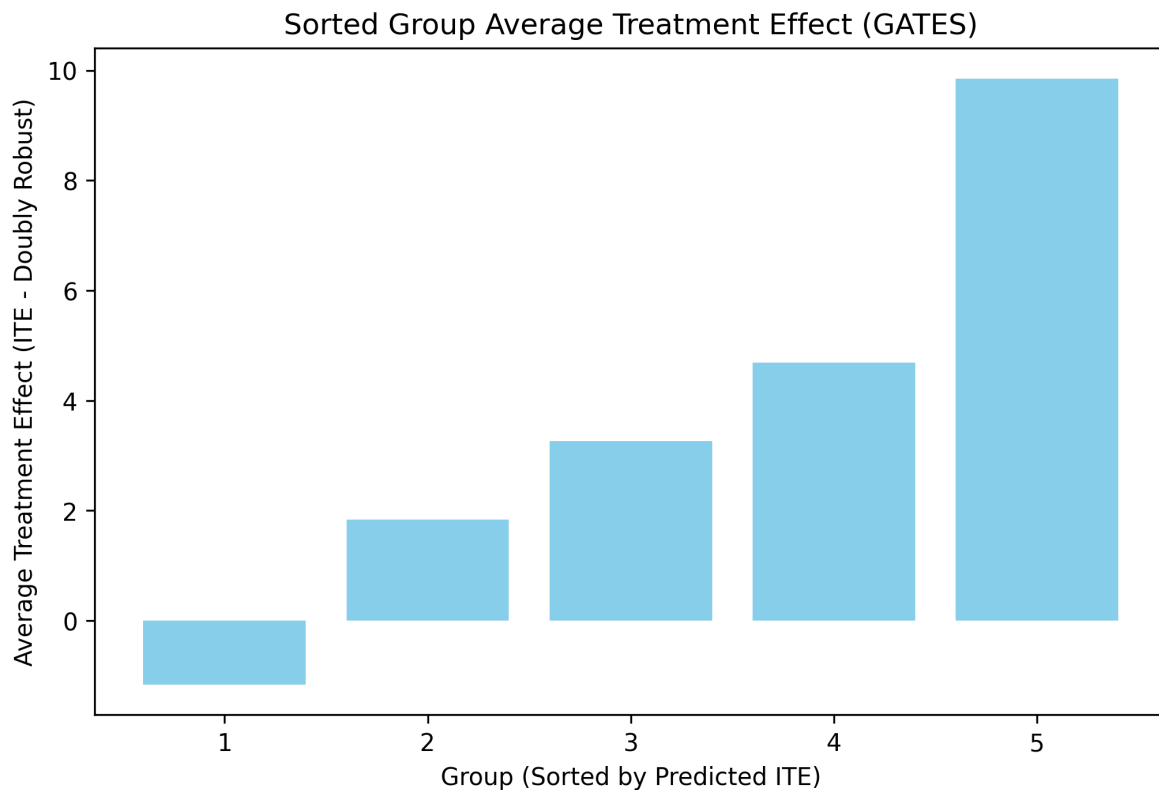
```

Sorted Group Average Treatment Effects (GATES):

```

group
1    -1.158407
2     1.832626
3     3.261341
4     4.693296
5     9.852490
Name: ite_dr, dtype: float64

```



## 7 Optimal Policy Learning

```
# Confusion matrix and report
cm = confusion_matrix(eval_data['optimal_treatment'], policy_preds)
print("Confusion Matrix:")
print(cm)

print("\nClassification Report:")
print(classification_report(eval_data['optimal_treatment'], policy_preds, zero_division=0))

# Plot the decision tree
plt.figure(figsize=(10, 7))
plot_tree(policy_tree, feature_names=X_eval.columns, class_names=["No", "Yes"], filled=True)
plt.title("Optimal Treatment Policy Tree")
plt.show()
```

```
[[ 165    64]
 [ 593 1178]]
```

	precision	recall	f1-score	support
0	0.22	0.72	0.33	229
1	0.95	0.67	0.78	1771
accuracy			0.67	2000
macro avg	0.58	0.69	0.56	2000
weighted avg	0.86	0.67	0.73	2000

[illegible]